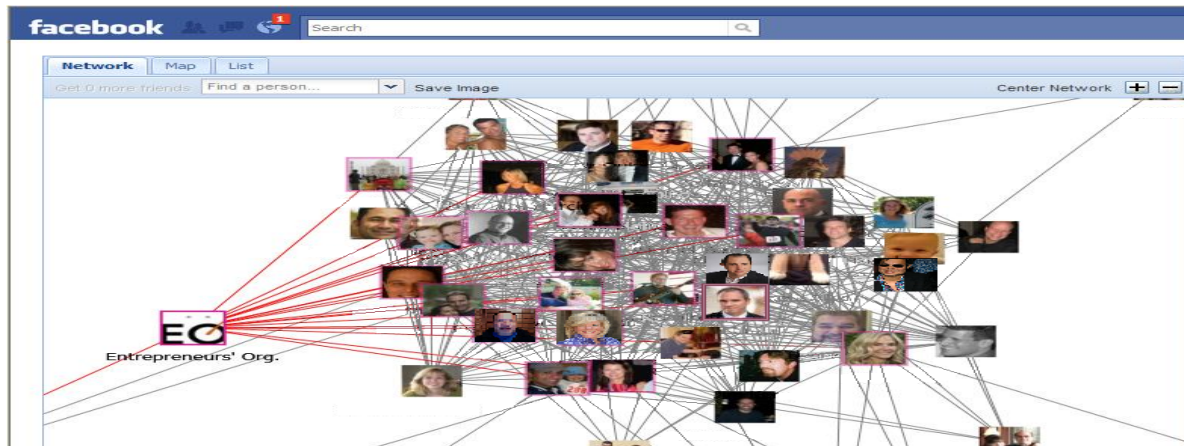


Problem

Networking & connecting with others is one of the most important features among many social media platforms. Facebook and other social media platforms need people to be able to connect users based on their common interests, degrees of separation, etc.



<http://www.fmsag.com/socialnetworkanalysis/facebook/>

Data

The dataset I used is from <https://snap.stanford.edu/data/ego-Facebook.html> with [Number] corresponding to file names: '0','107','348','414','686','698','1684','1912','3437','3980'. It contains the following files:

1. facebook_combined.txt -> Gives all of the connected node pairs

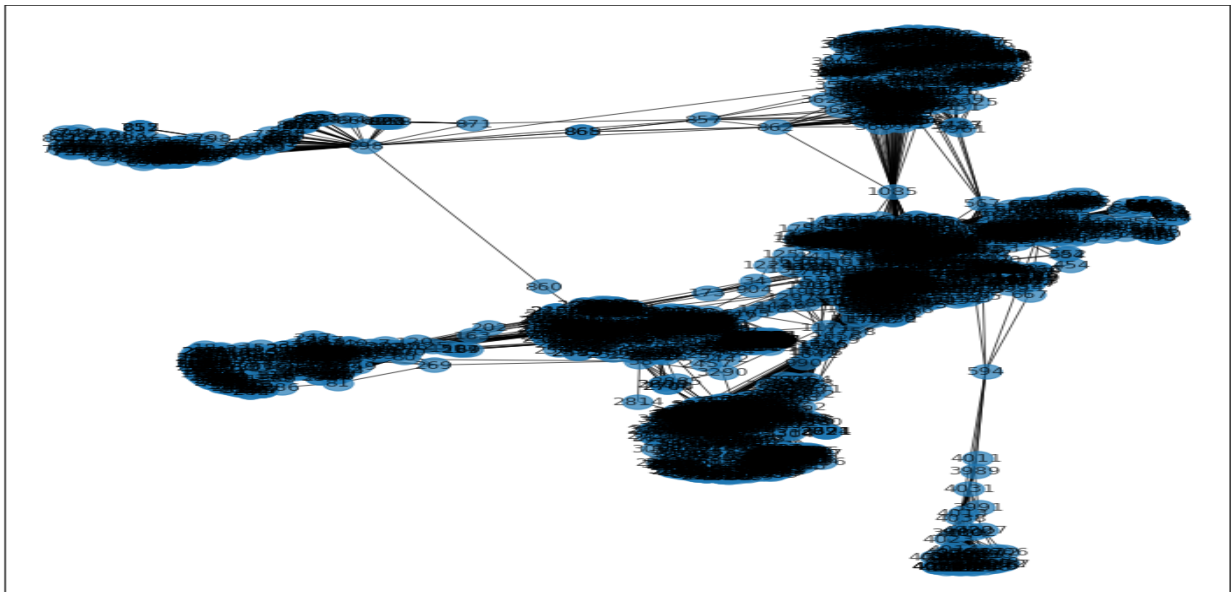
	Node 1	Node 2
	0	1
	1	2
	2	3
	3	4
	4	5
...
88229	4026	4030
88230	4027	4031
88231	4027	4032
88232	4027	4038
88233	4031	4038
88234 rows x 2 columns		

2. [Number].featnames -> Gives the feature type, semicolon, followed by specific feature value characterized by a number


```
{1: {'gender': [77], 'locale': [127]}, 2: {'education;school;id': [35], 'education;type': [53, 55], 'education;year;id': [57], 'gender': [78], 'languages;id': [92, 98], 'last_name': [114], 'locale': [126], 'location;id': [135]}, 3: {'birthday': [7], 'education;concentration;id': [14], 'education;school;id': [34, 50], 'education;type': [53, 55], 'education;year;id': [59, 65], 'gender': [78], 'languages;id': [92], 'locale': [127], 'location;id': [137], 'work;end_date': [168, 170], 'work;location;id': [137], 'work;start_date': [164, 202]}, 4: {'education;school;id': [50], 'education;type': [53, 55], 'education;with;id': [56], 'gender': [78], 'locale': [127]}, 5: {'education;school;id': [49, 50], 'education;type': [53, 54], 'education;year;id': [65], 'gender': [78], 'locale': [127]}, 6: {'birthday': [1], 'education;type': [53, 55], 'education;year;id': [62], 'gender': [78], 'last_name': [111], 'locale': [127], 'work;end_date': [157], 'work;start_date': [157]}, 7: {'education;concentration;id': [13], 'education;school;id': [25, 43, 50], 'education;type': [53, 54, 55], 'education;year;id': [59], 'gender': [78], 'last_name': [107], 'locale': [127], 'location;id': [137], 'work;employer;id': [141, 144], 'work;start_date': [196]}, 8: {'gender': [78], 'locale': [127]}}
```

Exploratory Data Analysis


I plotted an image of a spring layout of the network of people:



Pre-Processing

Node2Vec Pre-Processing

The goal was to generate features for the model to be able to use with a supervised learning model to predict if a pair of nodes are connected or not. The Node2Vec algorithm generates features for a node and needs to learn how to generate those features. To do this, Node2Vec needs to understand a previous version of the graph to generate the appropriate features for new connected or unconnected edges. Below shows a before and after images of deleting removable edges shown by the dotted lines:



The Node2Vec process works by Node2vec generating numerical representations of nodes in the graph via 2nd order (biased) random walk. First order random walk is done by sampling nodes on the graph along the edges of the graph, and each step depends only on the current state. Second order random walk is a modified version of the first order random walk that depends not only on the current state but also the previous state. A corpus of random walks is generated using each node in the network as a starting point. This corpus is then fed through word2vec to generate final node embeddings.

I followed the following steps in order to complete the Node2Vec pre-processing steps to generate features:

1. Found the negative samples (unconnected edges / node pairs) with nodes at max path length of 2 from one another to get samples that are possibly more likely to form a connection due to more common neighbors.

Unconnected Node Pairs: 1395922

	Node 1	Node 2	Connected
0	0	348	0
1	0	351	0
2	0	353	0
3	0	363	0
4	0	364	0
...
1395917	4035	4037	0
1395918	4035	4038	0
1395919	4036	4037	0
1395920	4036	4038	0
1395921	4037	4038	0

1395922 rows x 3 columns

2. Found the positive samples (removable edges / node pairs) - edges that can be removed while preserving the base graph structure (not eliminating nodes & not splitting the graph).

Removable Node Pairs: 84196

	Node 1	Node 2	Connected
0	0	1	1
1	0	2	1
2	0	3	1
3	0	4	1
4	0	5	1
...
88219	4020	4030	1
88220	4020	4031	1
88223	4021	4026	1
88226	4023	4031	1
88230	4027	4031	1

84196 rows x 3 columns

- Found the remaining edges by removing the positive edges from the list of connected edges so we can train the Node2Vec Model on those remaining edges (base structure of the graph).

Node2Vec Edges: 4038

	Node 1	Node 2
0	0	11
1	0	12
2	0	15
3	0	18
4	0	37
...
4033	4023	4038
4034	4026	4030
4035	4027	4032
4036	4027	4038
4037	4031	4038

4038 rows x 2 columns

- Combined negative & positive samples into final dataframe after reducing the # of negative samples to balance the final dataset to about equal positive & negative samples to be used for more accurate modeling.

Unconnected & Removable Edges: 169196

0 85000
1 84196
Name: Connected, dtype: int64

	Node 1	Node 2	Connected
0	0	364	0
1	0	906	0
2	0	961	0
3	0	970	0
4	0	978	0
...
169191	4020	4030	1
169192	4020	4031	1
169193	4021	4026	1
169194	4023	4031	1
169195	4027	4031	1

169196 rows x 3 columns

- Used the trained Node2Vec model to generate features for each node in a pair & sum them for a final set of features for that pair. Then, I test train split data for modeling.

```

Features (X): (169196, 100)
[[ 0.05423658  0.38881892 -0.7825276 ... -0.3608516  0.13439018
 -0.14539672]
 [-0.06878684  0.9884709  -0.85741556 ... -0.45005986 -0.20884675
 -0.87278837]
 [-0.4779037  0.9238887  -1.0211918 ... -0.7530286  0.5291232
 -0.5641365 ]
 ...
 [-0.29917532  0.92449725 -1.1229932 ... -0.36626023 -0.29428327
 -0.489528 ]
 [ 0.3246755  1.0931772  -0.41049516 ... -0.01053643 -0.36503953
 0.14392054]
 [ 0.48538733  0.91636264 -0.595041 ... 0.15409005 -0.65976167
 0.33824128]]

```

NetworkX Link Prediction Pre-Processing

I used the NetworkX library link prediction algorithms to calculate the link metric for each pair in the list of unconnected (negative samples) & removable edges (positive samples) to generate features that will be applied to a supervised learning approach to predict probabilities of each node pairs in forming a connection. I used the following link prediction algorithms: [common_neighbors](#), [jaccard_coefficient](#), [resource_allocation_index](#), [adamic_adar_index](#), & [preferential_attachment](#). Created & scaled the training set and testing set of data to create the final data below to be used for modeling:

	Node 1	Node 2	Connected	Common Neighbors	Jaccard Coefficient	Resource Allocation Index	Adamic Adar Index	Preferential Attachment
0	0	364	0	-0.688733	-0.940612	-0.639678	-0.704201	-0.371761
1	0	906	0	-0.688733	-0.942633	-0.871550	-0.733825	1.547445
2	0	961	0	-0.688733	-0.940750	-0.871550	-0.733825	-0.260503
3	0	970	0	-0.688733	-0.940818	-0.871550	-0.733825	-0.204874
4	0	978	0	-0.688733	-0.943233	-0.871550	-0.733825	2.270625
...
169191	4020	4030	1	-0.665866	-0.600397	-0.545057	-0.672454	-0.609909
169192	4020	4031	1	-0.597264	0.621218	0.571919	-0.502564	-0.615039
169193	4021	4026	1	-0.574396	0.936068	0.460927	-0.483467	-0.614157
169194	4023	4031	1	-0.597264	-0.034718	0.638937	-0.497499	-0.606222
169195	4027	4031	1	-0.620131	0.306369	0.240177	-0.556099	-0.615921

The link prediction algorithms employ the following equations to create the feature values and revolve around using the set of common neighbors between two nodes in a pair.

Modeling

I trained & tested 4 ML models to predict my binary classification of 1 or 0 representing connected status of connected or not connected respectively. I used logistic regression, Random Forest Classifier, Gradient Boosting Classifier, and a Multi-Layer Perceptron Classifier. I wanted a simple classifier so I picked a

logistic regression classifier, a random forest and gradient boosting classifier to test tree based classifiers, and the mlp classifier to try a neural network classifier to test my data on.


I took the following 5 steps to implement each of the 4 ML models:

1. I used GridSearchCV for hyperparameter tuning to pick the best model version.
2. I fit the model and made predictions for the test set.
3. Printed the classification report displaying the accuracy, precision, recall, and F1-scores.
4. Calculated the accuracy, precision, recall, F1-score, log loss score, and ROC-AUC (area under the curve) for the model.
5. Lastly, displayed the confusion matrix to see the distribution of predictions being made across all genres.

MLP Classification Report



	precision	recall	f1-score	support
0	0.95	0.88	0.91	17000
1	0.88	0.96	0.92	16840
accuracy			0.92	33840
macro avg	0.92	0.92	0.92	33840
weighted avg	0.92	0.92	0.92	33840



MLP Confusion Matrix




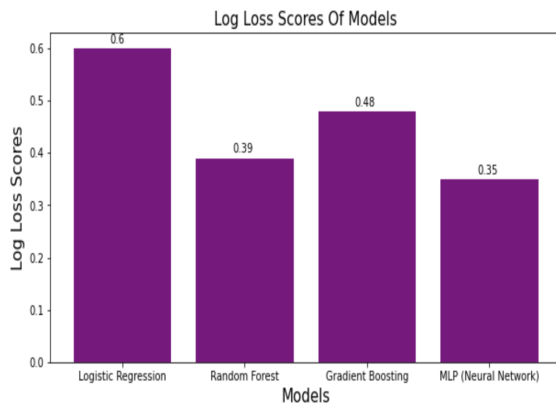
Node2Vec Model Comparison

The best model is the Multilayer Perceptron Classifier (Neural Network Classifier). The delta between the best and worst model in accuracy: 17%, precision: 18%, recall: 17%, f1: 17%. Here are the accuracies, precisions, recalls, and F1 scores for all 4 models.







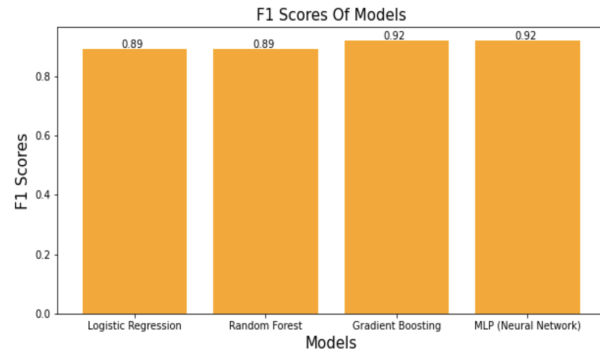

The lowest Log Loss score was 0.35 for the Multilayer Perceptron Classifier model. The ROC curves and the area under the curve (AUC) is the largest at 0.92 for the Multilayer Perceptron Classifier model.





NetworkX LP Model Comparison

The best models were the Multilayer Perceptron Classifier (Neural Network Classifier) & Gradient Boosting Classifier. The delta between the best and worst model in accuracy: 3%, precision: 2%, recall: 3%, f1: 3%. Here are the accuracies, precisions, recalls, and F1 scores for all 4 models:





The lowest Log Loss score was 0.24 for the Multilayer Perceptron Classifier & Gradient Boosting Classifier models. The ROC curves and the area under the curve (AUC) is the largest at 0.96 for the GB model and 0.95 for the Multilayer Perceptron Classifier model & the LR model.



Node2Vec & NetworkX LP Recommender Functions

I created a `Recommend_friends` function which gets the probabilities for each unconnected pair that includes the user requested user # and prints the top k as requested by the user. I took the following steps to create the functions:

1. Prompts to pick the user # to recommend friends for.
2. Prompts the user for # of recommendations.
3. Uses MLP classifier & 'predict_proba' method to get the probabilities of all unconnected node pairs being connected.
4. Ranks the probabilities and prints out the requested number of recommended friends user #s.

Node2Vec Function

```
In [339]: 1 recommend_friends()
Please enter the user (0-4038) to get friend recommendations for:
65
Please enter the k number of friend recommendations out of 22 that you
would like to receive:
22
Friend recommendations for user 65 are:
User 168 at 82.56% chance of a future link
User 221 at 73.05% chance of a future link
User 315 at 66.85% chance of a future link
User 56 at 62.53% chance of a future link
User 274 at 54.61% chance of a future link
User 75 at 53.54% chance of a future link
User 325 at 53.07% chance of a future link
User 3 at 48.2% chance of a future link
User 30 at 42.2% chance of a future link
User 73 at 29.01% chance of a future link
User 234 at 18.27% chance of a future link
User 229 at 17.61% chance of a future link
User 55 at 13.17% chance of a future link
User 92 at 10.66% chance of a future link
User 195 at 10.35% chance of a future link
User 177 at 4.59% chance of a future link
User 139 at 3.42% chance of a future link
User 167 at 0.55% chance of a future link
User 81 at 0.35% chance of a future link
User 93 at 0.32% chance of a future link
User 216 at 0.21% chance of a future link
User 6 at 0.03% chance of a future link
```

NetworkX LP Function

```
In [184]: 1 recommend_friends_nx()
Please enter the user (0-4038) to get friend recommendations for:
65
Please enter the k number of friend recommendations out of 22 that you would like to receive:
22
Friend recommendations for user 65 are:
User: 56.0 at 81.09% chance of a future link
User: 325.0 at 74.75% chance of a future link
User: 315.0 at 63.15% chance of a future link
User: 168.0 at 51.07% chance of a future link
User: 55.0 at 35.69% chance of a future link
User: 221.0 at 5.89% chance of a future link
User: 73.0 at 5.8% chance of a future link
User: 3.0 at 5.69% chance of a future link
User: 216.0 at 1.55% chance of a future link
User: 234.0 at 1.55% chance of a future link
User: 81.0 at 1.48% chance of a future link
User: 229.0 at 1.35% chance of a future link
User: 6.0 at 1.35% chance of a future link
User: 167.0 at 1.32% chance of a future link
User: 93.0 at 1.28% chance of a future link
User: 195.0 at 1.26% chance of a future link
User: 139.0 at 1.26% chance of a future link
User: 177.0 at 1.21% chance of a future link
User: 75.0 at 1.15% chance of a future link
User: 274.0 at 1.15% chance of a future link
User: 30.0 at 1.1% chance of a future link
User: 92.0 at 1.06% chance of a future link
```

Takeaways

In the Node2Vec & NetworkX features creation pre-processing approaches, the Multi-layer perceptron classifier was the best model & the gradient boosting classifier as well on the NetworkX LP approach. Also, the features generated by node2vec don't work for predicting links as well as the features generated by the NetworkX library like prediction algorithms.

Model	Hyperparameters	Accuracy	Precision	Recall	F1-Score	Log Loss	ROC-AUC
Node2Vec: MLP Classifier (Multi-layer Perceptron Classifier)	activation='tanh', max_iter=1000, random_state=0, solver='sgd'	0.85	0.86	0.85	0.85	0.35	0.92
NetworkX: MLP Classifier (Multi-layer Perceptron Classifier)	activation='tanh', max_iter=1000, random_state=0, solver='sgd'	0.92	0.92	0.92	0.92	0.24	0.95
NetworkX: Gradient Boosting Classifier	learning_rate=0.05, random_state=0	0.92	0.92	0.92	0.92	0.24	0.96

Further Research

I would try to implement different approaches to get link predictions other than the Similarity-Based Local Approaches (NetworkX LP) & Path & Walk-Based Method (Node2Vec).

